

PRODUCT REQUIREMENTS DOCUMENT

Seller Copilot

An autonomous growth agent for Mercado Livre marketplace sellers — from advice to action

Author: Felipe de Almeida · Strategy & Product

Type: Portfolio case study (1-pager PRD) · Status: Draft v1

Disclaimer: Independent strategy & product exercise. Not affiliated with or endorsed by Mercado Livre. Market figures are from public sources (see Sources) and used illustratively; all targets are hypothetical.

1. Context & opportunity

Mercado Livre is the largest e-commerce and fintech ecosystem in Latin America. Momentum is accelerating: Q4 2025 GMV reached US\$19.9B (+37% YoY), with Brazil and Mexico each growing ~35%; Q1 2026 GMV grew 38% YoY (FX-neutral) and items sold rose 56% YoY. Growth on this scale is driven by the supply side — the millions of third-party sellers who list, price, and fulfill the catalog.

Mercado Livre is already investing in AI on the seller side: in October 2025 it launched an AI Assistant that handled 9M+ conversations and a conversational Seller Assistant offering personalized advice. These tools answer questions well, but they are reactive: the seller must know what to ask and then act manually. The next unlock is shifting from answering to acting — closing the loop between insight and execution.

2. Problem statement

Most marketplace sellers are small operators without analysts or pricing teams. They are time-poor and react late to signals that directly affect sales:

- **Pricing drift:** they miss competitor price moves and Buy Box loss until sales drop.
- **Stockouts & overstock:** weak demand forecasting causes lost sales or frozen working capital.
- **Listing quality:** thin titles, attributes, and photos suppress search ranking and conversion.
- **Ad inefficiency:** ad budgets are set once and rarely optimized by ROAS.

Each gap is a known lever, but acting on all of them continuously is beyond a single operator. Advice alone does not close it; execution does.

3. Proposed solution

Seller Copilot is a proactive agent layered on top of the existing Seller Assistant. It continuously monitors each seller's catalog and market, surfaces the highest-impact actions, explains the expected outcome, and — with the seller's permission — executes them. The seller stays in control through three autonomy levels.

Autonomy level	Copilot behavior	Best for
1 · Advisory	Recommends actions with projected impact; seller executes manually.	New / cautious sellers

Autonomy level	Copilot behavior	Best for
2 · Approve-to-act	Prepares the change; executes on one-tap approval.	Most sellers
3 · Autopilot	Executes within seller-set guardrails (e.g. price floor, max ad spend); reports after.	High-volume sellers

3.1 MVP scope (first three skills)

- **Dynamic pricing:** monitors competitor prices and Buy Box status; recommends/executes price moves within a seller-defined floor and margin rules.
- **Inventory radar:** forecasts demand per SKU; flags stockout and overstock risk with reorder suggestions.
- **Listing optimizer:** uses generative AI to rewrite titles, fill missing attributes, and flag weak images, with before/after ranking estimates.

Explicitly out of MVP: ad-spend automation, multi-channel sync, and supplier negotiation (see roadmap).

4. Success metrics

North Star: incremental GMV per active seller attributable to Copilot actions (measured against a holdout control group).

Metric	Type	Target (90 days post-launch)
Incremental GMV / active seller	North Star	+8% vs. holdout
Copilot-assisted sellers (adoption)	Input	25% of eligible sellers
Recommendation acceptance rate	Engagement	≥ 40%
Buy Box win rate (pricing skill)	Outcome	+10 pts on enrolled SKUs
Stockout days / seller	Outcome	-20%
Action reversal / complaint rate	Trust guardrail	< 3%

5. Key trade-offs & risks

Decision / risk	Tension	Position
Autonomy vs. trust	Acting for sellers risks errors that erode trust.	Ship at Level 2 (approve-to-act) by default; Autopilot is opt-in with hard guardrails.
Model cost/quality/latency	Best models are costly; sellers expect instant, cheap.	Route by task: small models for monitoring/classification, larger models only for generation; cache aggressively.
Hallucination	Wrong price or claim has real financial impact.	Constrain actions to validated tools and rules; no free-text actions; every action is auditable & reversible.
Build vs. extend	New surface vs. existing Seller	Extend the assistant; reuse its

Decision / risk	Tension	Position
	Assistant.	context and trust, add an action layer.

6. Phased roadmap

Phase	Timeline	Focus
0 · Discovery	Weeks 0–4	Seller interviews, data audit, holdout design, guardrail rules.
1 · MVP (Advisory + Approve)	Weeks 5–16	Pricing + Inventory + Listing skills at autonomy levels 1–2; closed beta.
2 · Autopilot + scale	Weeks 17–28	Level-3 guardrails, broader rollout, ad-spend skill.
3 · Ecosystem	Weeks 29+	Mercado Pago credit signals, cross-border, supplier negotiation.

7. Open questions

- Pricing model: free growth lever, or premium tier / take-rate on incremental GMV?
- How do we attribute incremental GMV cleanly without contaminating the holdout?
- What is the minimum guardrail set required for Autopilot to be safe and compliant?

Sources

- Mercado Libre Q4 2025 sales & GMV growth — <https://mexicobusiness.news/ecommerce/news/mercado-libre-delivers-record-2025-sales-45-4q25-growth>
- Mercado Libre Q1 2026 results (GMV & items sold) — <https://www.businesswire.com/news/home/20260507238255/en/Mercado-Libre-Kicks-Off-2026-with-Fastest-Revenue-Growth-in-Almost-Four-Years-as-Strategic-Investments-Drive-Market-Share-Gains>
- MercadoLibre AI investment & GMV growth in Brazil/Mexico — <https://seekingalpha.com/news/4556373-mercadolibre-outlines-continued-investment-in-ai-and-shipping-as-gmv-grows-35-percent-in>